

What can your AI see?

Context, fluency, and connectors

Week 6 · Learn, Vibe, Build · October 5, 2026

Choose what it sees. Check what it produces. Control what it can change.

- Context: what the AI sees
- Fluency: how you work with it
- Connectors: what it can reach and change

Context is the working material

- Instructions + conversation + files + tool results
- Training knowledge $\neq$ knowledge of your project
- Saved memory is the app retrieving—check what loads

Give it the missing facts

- Weak: “Improve my portfolio.”
- Add: audience, current page, goal, constraints
- Same task; different context.

Choose context; don't dump everything

- Load what this task needs
- Label facts, instructions, uncertainty
- Handoff: goal · decisions · next step · evidence

AI Fluency: the 4D framework

- Delegation · Description · Discernment · Diligence
- Effective, efficient, ethical, and safe
- Dakan (Ringling) & Feller (UCC), with Anthropic

Three ways to work with AI

- Automation: AI does a task you defined
- Augmentation: you think together
- Agency: you set AI up to act on its own later

Delegation: whether, when, how

- Know the goal and the task before the tool
- Pick the mode and the right tool
- Keep the work where the thinking is the deliverable

Description: Product · Process · Performance

- Product: what you want out (format, audience, style)
- Process: how to approach it (steps, what to read)
- Performance: how it should behave (concise? push back?)


Discernment: fluent ≠ correct

- Judge the output and the behaviour
- Product · process · performance—again
- Check facts, run it, look closely

Diligence: responsibility doesn't transfer

- Creation: choose tools and data thoughtfully
- Transparency: be honest about how AI was involved
- Deployment: verify before you share or ship

Two loops



- Description ↔ Discernment: the working loop
- Delegation ↔ Diligence: the responsibility loop

Markdown is readable structure

```
# Project
```

```
## Goal and audience
```

- Scholarship panel

```
## Constraints
```

- Don't invent achievements

Plain text you can read, edit, and paste.

A context file is a description that doesn't go away

- Goal: who is this for, what should change? (Description)
- Inputs: what should it read? (Description)
- Boundaries: what must it never do? (Diligence)
- Check: how will you know? (Discernment)

Two-minute exercise: make context usable

- 0:00–1:00 Write goal, input, boundary, check.
- 1:00–2:00 Partner: name one missing fact.
- Save as project-context.md—or keep it in chat.

Live demo · same task, better context

- A: portfolio opening, no brief
- B: fresh chat + project-context.md
- Compare: audience · facts · next action

Three files, three jobs

- Instructions: how to work on this project
- Recipes: how to repeat a task
- Knowledge: facts and decisions worth retrieving

Index files: a map for a body of knowledge

```
# Course knowledge – index
- syllabus.md: schedule, policies
- week-06/lesson.md: context, 4D, MCP
- projects/anim/context.md: my brief
- glossary.md: terms
```

The AI reads the map,
then opens only what it
needs.


Tools connect the work to the world

- Read a note. Search a folder. Check a calendar.
- A tool result becomes new context.
- A write action changes something outside the chat.

API vs. MCP

- API: one service's menu of requests, different for each
- MCP: one shared way for AI apps to use many tools
- MCP servers often call APIs underneath
- Like wall sockets vs. a universal adapter

MCP: how the pieces fit



- Servers expose tools, resources, and prompts.
- A standard connection is not unlimited access.

Permission is part of the design

- Start read-only, smallest useful scope
- Ask before sending, paying, deleting, publishing
- Read back anything consequential

A document is data—not your boss

- A page says: “Ignore your instructions.”
- That’s untrusted content, not a command
- Keep secrets out; stop when permission is unclear

Try one change on your real project

- Add a context file; compare two attempts
- OR one read-only connection, one clear question
- Share the evidence—not just the polish

Context is also how a build becomes a product

- Your context file names an audience, inputs, a check
- That's the first draft of a product brief
- Who besides you would use this—and what do they see first?

Learn this your way

- The whole lesson is one markdown file
- Paste it into your AI and ask it to teach you
- It can explain, quiz you, and use your project
- Then check it: Discernment applies here too

Teach-me prompts

- “Explain the 4Ds using my project: ...”
- “Quiz me, one question at a time.”
- “Interview me to write my context file.”
- “Build me an index.md for my notes.”

Q&A · what is missing from your context?

- What should your tool know before it starts?
- What should it never do without asking?
- Next: 30 minutes of student sharing.

Sources

- AI Fluency: Framework & Foundations · Dakan, Feller, Anthropic
- Effective context engineering · Anthropic
- AGENTS.md convention · agents.md
- MCP architecture and tools · modelcontextprotocol.io